ASHITA CHANDNANI

Boise, ID (relocating to San Jose, CA, 2027) | ashita0904@gmail.com | linkedin.com/in/ashita-chandnani | github.com/AshitaC | ashitac.github.io

SUMMARY

MS in Artificial Intelligence student at UT Austin with an M.S. in Computer Science (4.0 GPA). I build LLM applications that work outside a notebook: agents with memory, retrieval systems that cite their sources, and ML pipelines that deploy through CI/CD on AWS. Research background in recommender systems and sensors. Seeking 2027 AI/ML engineering internships and roles.

EDUCATION

M.S., Artificial Intelligence | *The University of Texas at Austin* — Expected Dec 2027

M.S., Computer Science | *Boise State University* — Aug 2025 | GPA 4.0/4.0

M.Tech., Electrical Engineering | *Indian Institute of Technology Gandhinagar* — CPI 9.2/10

B.E., Electronics & Communication Engineering | *Govt. Engineering College Bikaner (Univ. of Rajasthan)* — First Class Honours

SKILLS

LLMs & agents: LangChain, LangGraph, ReAct agents, RAG, prompt engineering, structured JSON output, Groq, Llama 3.3, OpenAI API
Machine learning: PyTorch (attention models), BERT embeddings, scikit-learn, TensorFlow, model evaluation (MRR, nDCG, Hit Rate)
Cloud & MLOps: AWS (Bedrock AgentCore, EC2, S3, ECR, CloudWatch), Docker, GitHub Actions CI/CD, FastAPI, Streamlit, MongoDB Atlas
Data: Python, SQL, PySpark, Spark SQL, Hadoop (HDFS), Pandas, NumPy, FAISS, ChromaDB, Power BI

PROJECTS

Conversational Agent with Long-Term Memory | *AWS Bedrock AgentCore, LangGraph, FAISS, Groq*

- Built and deployed a stateful FAQ agent on AWS Bedrock AgentCore with session memory and cross-session long-term memory.
- Designed a LangGraph ReAct agent with three retrieval tools over a FAISS vector store; memory middleware injects relevant past context before each LLM call.

Semiconductor Research Assistant (RAG) | *LangChain, ChromaDB, Llama 3.3, Streamlit*

- Built and deployed a RAG assistant that answers semiconductor research questions with citations from uploaded PDFs and URLs, with session-isolated storage.

MLOps Pipeline for Insurance Lead Prediction | *FastAPI, Docker, AWS, GitHub Actions, MongoDB Atlas*

- Built an end-to-end pipeline from data ingestion to deployment with modular validation, transformation and training components.
- Automated CI/CD with GitHub Actions and Docker, deploying to AWS EC2 via ECR; S3 as model registry.

Financial KPI Extractor | *LangChain, Llama 3.3, Pydantic, Streamlit*

- Extracted revenue and EPS from unstructured financial news into schema-validated JSON, with Plotly charts and CSV export.

MovieLens Data Mining at Scale | *PySpark, Spark SQL, Hadoop (HDFS)*

- Answered 10+ analytical questions on 1M ratings with Spark SQL; cluster run finished in 46 seconds versus hours on a single VM.

RESEARCH EXPERIENCE

Graduate Researcher, Computer Science | *Boise State University* — Aug 2024 – Aug 2025

- Re-implemented FANAR, a trust-aware news recommender, in PyTorch: attention over frozen BERT embeddings with a Beta trust model.
- Achieved **48% higher MRR@10** than the strongest of 7 baselines on PolitiFact; trust filtering alone added 18.8% MRR@10. Paper under review at ECIR 2027.

Graduate Researcher, Electrical & Computer Engineering (Doctoral Program) | *Boise State University* — 2017 – Dec 2020

- Printed graphene and carbon-nanotube sensors (IGEM fellowship). Equal-contribution first author, Sensors review (2020, 120+ citations); co-inventor, US 2018/0342976 A1.

Assistant Professor, Electronics & Communication Engineering | *Nirma University, Ahmedabad* — 2014 – 2015

- Taught Digital Electronics and Communication Systems; mentored undergraduate student projects.

LEADERSHIP & INITIATIVE

AI Bites: started and write a LinkedIn series explaining AI concepts in plain language; one post reached 68,000+ impressions.
Certifications: Linear Algebra and Calculus for ML (DeepLearning.AI); Statistics (Stanford Online); Data Science with Python (Edureka); SQL, Python, Power BI (Codebasics).